

X (Twitter) — community, distribution and revenue

Status: strategy + build plan. Reconnaissance complete, transport absent. **Requested by:** owner, 2026-09-15. **Supersedes nothing.** `docs/product/twitter-bot-voice-spec.md` remains the voice authority for what the bot says; this document is about what X is *for* and what is missing.

1. What already exists (so this does not get rebuilt)

Asset	Path	State
Voice spec — four event kinds, eight hard refusals	<code>docs/product/twitter-bot-voice-spec.md</code>	Complete
Templates	<code>apps/web/lib/twitter-bot/templates/*.ts</code>	Built — two doctrine conflicts, see §3
Planner — idempotency, blocked reasons, compliance gate	<code>apps/web/lib/bot-outbox/plan.ts</code>	Built
DB adapters	<code>apps/web/lib/bot-outbox/records.ts</code>	Built
Operator review UI	<code>apps/web/app/cockpit/bot-outbox/page.tsx</code>	Built
Eval cases	<code>docs/ops/evals/twitter-bot-*.md</code>	Written
Transport — auth, rate limit, send, retry, log	<code>workers/twitter-bot/</code>	ABSENT
Community / engagement / partnership layer	—	ABSENT

The finding, stated plainly: GSE has a fully specified, compliance-gated, operator-reviewable X bot that cannot post. Everything up to the HTTP call exists. That is an unusually well-defined gap, which is why this is cheap to close.

2. What X is for — four jobs, in priority order

The account is `@GalaxySportsAI`. Four jobs, and they are not equal. Getting the order wrong is how a trust product becomes a content farm.

Job 1 — The public record needs a witness (highest value) The product’s entire claim is that its record is checkable. Today the record is checkable *by whoever visits the site*. X changes that: a settlement posted at the moment it settles is timestamped by a third party, publicly, before the outcome can be curated out.

This is the single most defensible thing the account can do, and it is also the hardest to fake. **Post every settled free-tier pick, wins and losses alike, with no filter.** A bot that posts 73% of its wins and 100% of its losses is a bot nobody can accuse of anything, which is the point.

Job 2 — Distribution (medium value, real) Every post links to a surface. Not “link in bio”, not a thread that ends in a funnel — a direct link to the specific `/room/[gameId]` or `/picks` page the post is about. The measurable goal is referrals that convert to a first repeat visit, not follower count.

Job 3 — Community (the owner’s ask — and the one to be most careful with)

There is a real appetite on X for a *methodology* community: people who want to talk about calibration, closing line value, sample size and variance rather than locks. That audience exists and is currently underserved — it is the same audience that reads `docs/calibration-proposals/`.

What that community looks like when it is on-brand: - A weekly published calibration post — the curve, the Brier score, and the verdict *including when the verdict is “higher confidence is winning less.”* - Loss autopsy threads. These are the highest-engagement content the account can produce, and they are also the most honest. Nobody else publishes them. - Open methodology questions answered with the framework and never the weights (`docs/positioning.md` § Public Methodology Rule already draws that line). - A model-changelog post on every `MODEL_VERSION` bump, naming what changed and why.

What it must never become (each is a hard refusal already in the voice spec, listed here because community-building pressure pushes toward all four): 1. “Who do you have tonight?” — engagement bait. 2. Reply-guy behaviour on other people’s posts — off-mission and dilutes the account. 3. Giving out picks in replies to people who ask. The bot has no ability to publish a pick that did not clear the gate, and replying to a request is exactly that. 4. A follower-count milestone post. Nobody who matters is impressed, and it is the “more action” impulse the strategy triage rejects.

Job 4 — Research input (real, but gated) X is the fastest surface for injury news, lineup changes and beat-reporter information. That is genuinely valuable for a market-reading model. It is also the fastest surface for rumours.

The repo already has the right primitive:

`apps/web/lib/decision-genome/rumor-quarantine.ts`. Any X-sourced signal goes through it or does not enter. The lane is specified in §5 and is deliberately last.

3. Two doctrine conflicts that must be fixed before the bot posts

Found during reconnaissance. Both are in the *existing* templates, so they are live defects the moment the transport ships.

Conflict 1 — the spec renders confidence as a percent `pick-publication.ts` produces:

That spec predates the 2026-09-13 measurement (AGENTS.md): confidence is a weighted factor sum whose 80+ band claims 0.8663 and realizes **0.5191** ($z = -10.7$). A percent on the *widest-reach surface the product owns* asserts a win probability the number demonstrably is not — and unlike the app, a tweet cannot be recalled.

Fix: render 72/100 and say “score”. Correct the spec and the template together, and add an eval.

Conflict 2 — the spec permits //⚖️ the design contract forbids emoji The voice spec calls them “the only emojis the bot uses”. DESIGN.md mandates W/L/P/V monograms and says “emoji $\approx$ zero”; apps/web/lib/positioning-vocab.json plus scripts/guardrails/trust-gate.mjs enforce the ban. Both cannot be right.

Resolution — a documented, single-use exception. Settlement glyphs are permitted **only on the settlement lead post**, and nowhere else, for a platform-specific reason: at timeline scale a monogram is illegible and the platform’s own convention for outcome carries meaning that a letter does not. Every other post uses the sanctioned data glyphs ($\uparrow \downarrow - \cdot \rightarrow$). The exception is written into the voice spec with its boundary, and an eval asserts a glyph does not appear on any other event kind.

4. Revenue on X — an honest reading

Option	Verdict	Reasoning
Sportsbook affiliate links	Reject	<code>platform-gaps-triage.md</code> §  is explicit: co-branded sportsbook checkout “turns us into a tout funnel that pushes wagering”. An honest, clearly-labelled disclosure link is the <i>most</i> that would ever be considered, and it is founder-gated.
Sponsored posts	Reject	The account’s value is that it has no reason to lie. Sponsorship prices exactly that.
X Premium / subscriptions revenue share	Not a business	Payouts are trivial at any realistic follower count, and the programme requires a paid tier.
Driving app subscriptions (IAP)	The actual revenue path	X is top-of-funnel for the App Store listing and the web pricing page. Measured as: link clicks → store page views → trials.
Creator/contributor programme (multi-contributor House picks)	Founder-gated, and phase-appropriate	<code>docs/product/phase-6-plus-planning.md</code> Item 2 specs it. It requires a solo creator layer that has run 6+ months first, plus a revenue-split decision. Not a v1 build, and building the tooling now would be building for a decision that has not been made.
B2B / data licensing visibility	Real, second-order	The same posts that demonstrate methodology rigour are the ones an institutional buyer reads. No product work; just do not hide the record.

The honest summary: X is a distribution and trust channel for this product, not a revenue channel. Any X revenue model that works is one where the account's credibility is the asset — which means the model must never be the thing that spends it.

5. The research lane — X as a signal source

Safe only with a strict pipeline, and the pipeline is the deliverable:

```
X post (beat reporter, injury news)
  → decision-genome/rumor-quarantine.ts    (existing; quarantines
unverifiable claims)
  → source-rights registry                  (existing; what may be read and
how used)
  → corroboration requirement: >= 2 independent sources OR an official team
statement
  → THEN, and only then, a signal category for the engine
```

Three rules that make this safe:

1. **No X-sourced signal may be the sole basis for a published pick.** Corroboration is required. A single reporter's claim moves a line, and moving lines are exactly what this model reads — so an uncorroborated rumour would become a self-fulfilling input.
 2. **The timestamp is the signal, not the claim.** "This was reported 40 minutes before the line moved" is measurable and useful. "This report is true" is not something X can tell us.
 3. **Nothing X-sourced is ever quoted.** The engine consumes a signal category, never text, both for rights reasons and because a rumour rendered as prose is a rumour the product now appears to endorse.
-

6. What to build, and in what order

#	Build	Why in this order
1	Transport — OAuth 1.0a signer, API v2 client, rate limiter, idempotent sender, <code>MUTE_BOT</code> honouring, attempt logging	Without it nothing posts and every other item is theory.
2	Fix the two template conflicts + evals	Must land <i>with</i> the transport, or the first post ships the error.
3	Wire the outbox to a scheduler (Vercel cron, matching the repo's existing 22 schedules)	The planner exists; it needs a clock.
4	Community cadence — weekly calibration post, autopsy thread, model changelog	Turns the account from a settlement log into a surface people follow.
5	Measurement — referral attribution from post → store page → trial	Without it, “did X work” is a feeling.
6	Research lane , through rumor-quarantine	Last, because it is the only item that can damage the model rather than the account.

Transport specification

- **Auth:** OAuth 1.0a user context (HMAC-SHA256). A signer is pure, deterministic, and unit-testable on any host — no network needed to verify the signature algorithm against RFC 5849 test vectors.
- **Endpoint:** `POST /2/tweets`, with `POST /2/tweets` + `reply.in_reply_to_tweet_id` for threads.
- **Rate limits:** honour `x-rate-limit-remaining` / `x-rate-limit-reset` from the response headers rather than a hardcoded ceiling; the account's tier determines the real limit and a stale constant is how a bot gets muted by the platform.
- **Idempotency:** the planner already emits an `idempotencyKey`. The transport records it BEFORE sending and checks it before sending, because the failure mode is “posted twice after a timeout”, not “posted never”.

- **MUTE_BOT** is checked at send time, not at schedule time. A mute that only takes effect on the next cycle is not a mute.
 - **Logging:** one row per attempt with outcome, so **AgentRunLog** answers “what did it post”.
 - **Account hygiene from the spec is automated, not aspirational:** the bot follows nothing, likes nothing, and retweets nothing. A test asserts the transport exposes no such call.
-

7. Metrics that matter (and the ones that do not)

Matter: referral clicks to `/room/*`, store page views attributable to X, trials started, first repeat visit rate, and the ratio of settlement posts to publications (a healthy feed is roughly 1:1 over a season — a feed that is mostly wins means losses are being suppressed).

Do not matter: follower count, impressions, likes. Each of them is trivially gameable by engagement bait, which is the one thing the voice spec forbids outright. A metric you are not allowed to optimise for is not a metric.